

A FIELD GUIDE • IRREVERENT MARKETING

Steal the *Signal*

*The DTC operator's field guide to
reverse-engineering any brand's Meta
ads — without the tools they think they
need.*

EDITION

OPERATOR'S FIELD GUIDE • 01

EIGHT MODULES • ONE KIT

M 00-07

Most competitor research is theater. This isn't that.

You can scroll a competitor's ad library for an hour and walk away with nothing. Or you can spend twenty minutes with the method in this book and walk away with a scored swipe vault that tells you exactly which of their ads is carrying the account — and why yours isn't.

WHO THIS IS FOR

DTC operators running \$10M–\$65M brands. In-house marketing leads sitting in meetings they're afraid of. Agency strategists whose pitch decks are built on vibes.

WHAT IT REPLACES

The guess. The vibe-check. The "our creative just feels tired." The three-hour scroll through a competitor's library that ends with a Slack screenshot and no plan.

WHAT YOU'LL BUILD

A populated Swipe Vault. A scoring workflow. A brief template that converts stolen signal into creative direction your team can execute on Monday.

HOW LONG IT TAKES

Ninety minutes to read. One afternoon to build the vault for three competitors. From there it's a weekly forty-five-minute rhythm that compounds.

The Eight Modules

Each module ends with a Do this now action. No filler. No screenshots of things you already know how to do.

M • 00	The Visibility Problem	p . 05
Why creative keeps dying on you. What the Meta Ad Library actually shows. The duration insight that replaces every metric you can't see.		
M • 01	The No-Tool Path	p . 13
Scraping any brand's ads with nothing but a browser and Claude. The exact prompt. What to capture and what to skip.		
M • 02	The Apify Path	p . 21
Scaling the scrape to five to fifty brands with one actor. Cost per run. The JSON you actually want. What to ignore in the output.		
M • 03	The Swipe Vault Database	p . 29
The schema. The hook taxonomy. Templates for Notion and Google Sheets — pick your poison, both work.		
M • 04	The Scoring Layer	p . 37
The duration ladder. Variant-count signals. Angle clustering. The winner-ladder framework that tells you what's actually working.		

M·05

Steal the Funnel, Not the Ad

p . 45

Where mid-market brands bleed. The ad-to-page mismatch audit. Why a winning ad doesn't save a broken offer.

M·06

Signal to Brief

p . 53

Converting stolen data into creative direction. The structural-DNA method. The brief template your studio or AI pipeline can execute verbatim.

M·07

What's Next

p . 61

The Five-Library Done-For-You. Upgrade paths into BIS and the Dashboard. The weekly rhythm that keeps you off the plateau.

HOW TO READ THIS

Skip around if you want. But if your swipe vault is empty, start at M·00. Every module after that assumes the one before it is done. The compounding is the point.

MODULE • 00 • THE REFRAME

The Visibility Problem.

You don't have a creative problem. You have a visibility problem. Once you can see what they're doing, you can't unsee it — and you can't go back to guessing.

READ TIME

18 minutes

ACTION AT END

Pick three competitors

PREREQUISITES

None

§ 01 · THE PLATEAU PATTERN

Your creative keeps dying and you can't see why.

The ad that did 3x ROAS last quarter is a cold corpse this quarter. You don't know why it died. You don't know what killed it. You don't know which of your three new tests is going to replace it — or whether any of them will.

This is the plateau. Every DTC brand between eight and sixty-five million in revenue hits it. The brand that was compounding stops compounding. Agency gets blamed. Creative team gets blamed. Media buyer gets blamed. Nobody inside the building can actually point to what changed because nobody has the visibility to see it.

The operators who break through the plateau don't have better creative teams. They don't have bigger budgets. They have **better visibility into what's working — their own, and everyone else's.**

The reason this course exists is simple. Meta's Ad Library — the free, public-facing database of every active ad on the platform — is the most under-used competitive intelligence tool in the direct-to-consumer stack. Operators know it exists. They've browsed it. A small fraction have a repeatable method for pulling signal out of it. Most don't.

That's the gap this course closes. By the end, you'll have a method that scales from one brand to fifty, and a database that compounds every week you run it.

§ 02 · THE TWO-COLUMN TRUTH

What the Library gives you. And what it refuses to.

The first thing to understand about the Ad Library is that it was built for political transparency, not marketing intelligence. This is why the disclosure layer is lopsided — and why most operators misread what's possible.

Here's the honest split.

VISIBLE · PUBLIC DISCLOSURE

What the Library hands you on every active ad

CREATIVE ASSET	Every image, video, carousel, collection. Full resolution.
AD COPY	Primary text, headline, description, CTA button, link preview.
LANDING PAGE	Destination URL. Follow it. This is where half the signal lives.
PLATFORMS	FB feed, IG feed, IG Stories, Reels, Messenger, Audience Network.
ACTIVE SINCE	Start date on every ad. This is the most valuable field on the page.
VARIANT COUNT	How many versions of the same concept are running in parallel.
PAGE CONTEXT	Follower count, verification, country of origin, creation date.

HIDDEN • PRIVATE TO THE ADVERTISER

What the Library will never show you (for US commercial brands)

SPEND	NOT DISCLOSED
-------	---------------

IMPRESSIONS	NOT DISCLOSED
-------------	---------------

REACH	NOT DISCLOSED
-------	---------------

FREQUENCY	NOT DISCLOSED
-----------	---------------

CTR / CPC / CPM	NOT DISCLOSED
-----------------	---------------

TARGETING	NOT DISCLOSED
-----------	---------------

DEMOGRAPHICS	NOT DISCLOSED
--------------	---------------

EXCEPTION • EU / EEA TARGETING

Ads delivered to European users disclose estimated reach ranges, country-level targeting, gender and age splits, and the parameters used to build the audience. If you're spying on a brand with European distribution, you get dramatically more. This is the Digital Services Act doing work Meta wouldn't do on its own.

§ 03 · THE DURATION INSIGHT

Brands don't burn budget on losers. This is your unlock.

Here's the reframe. Every operator's first instinct when they see the hidden column is: "Then how am I supposed to know what's working?" The answer is already on the page, and most people miss it.

An ad that's been active for forty-five days is not surviving by accident. It's surviving because it's the ad paying the bills.

THE CORE REFRAME · M · 00 · 03

No media buyer in the world leaves a losing ad running for six weeks. Budgets don't allow for it. Account performance doesn't allow for it. The CFO doesn't allow for it. If an ad is still live thirty days in, it's paying for itself — and the longer it's been alive, the harder it's paying.

Duration is your proxy for spend. Not a perfect one. But a directionally honest one, and the only one the Library will ever give you.

The reverse is also true. An ad that went live and died inside a week is almost certainly a loser — killed by the brand's own performance thresholds. You don't need impressions data to know this. You need a start date.

The duration ladder

> 30 days	Control winner	STEAL THIS
7 - 30 days	In testing	WATCH IT
< 7 days	Fresh test	UNPROVEN
Killed < 7d	Dead on arrival	IGNORE

Layer variant count on top of the ladder and the picture sharpens further. A single long-running ad with no variants is a brand that cracked a hook and is milking it. A long-running concept with twelve live variants is a brand that cracked the angle and is iterating the creative to delay fatigue. Both are winners. The variant count tells you how mature their testing is.

Inversely: a brand running thirty new ads a week with nothing older than fourteen days is a brand that's thrashing. They don't have a winner. Their account is noise. You don't need to copy anything they're doing — you need to avoid what they're doing.

§ 04 · WHY THIS COURSE EXISTS

You already know the Library is there. That's not the gap.

Most operators have opened the Ad Library. They've typed a competitor name in. They've scrolled. And then they've closed the tab, because there is no obvious method for converting two hundred active ads into a decision they can make on Monday morning.

The gap isn't access. The gap is **method**. This course is the method. Five steps, repeatable:

- **Scrape** — pull the data, with or without tools. Both paths taught.

- **Structure** — land the data in a schema you can query. Swipe Vault.
- **Score** — apply the duration ladder and hook taxonomy. Rank.
- **Steal** — extract the structural DNA. Not the creative. The pattern.
- **Steer** — convert the pattern into a brief. Ship.

Every module that follows is one of those five verbs, broken into an operator's playbook with an action at the end. By the time you finish M-06, you'll be able to read any DTC brand's ad account without seeing a single internal metric — and know, within ten minutes, what's carrying it.

§ 05 · WHAT YOU'LL HAVE WHEN YOU FINISH

The deliverables, stated plainly.

A populated Swipe Vault. Your own database of fifty or more scored competitor ads, tagged by hook type, ranked by duration, and cross-referenced against landing page offer match. Living document — you'll add to it weekly.

A repeatable scraping workflow. Two paths, pick one. The no-tool Claude path for operators who don't want another SaaS subscription. The Apify path for operators scaling past five brands at a time.

A scoring framework. The duration ladder plus the hook taxonomy plus the funnel-match audit. Three lenses. A winning ad has to pass all three. Losers fail at least one.

A brief template. The structural-DNA fill-in. It converts a scored competitor ad into a brief your creative team — or your AI pipeline — can produce work from without copying anyone.

A weekly rhythm. Forty-five minutes, once a week, to keep the vault current and the signal fresh. The compounding is where the real value lives. Most operators stop at the one-time scrape. That's a mistake.

Optional: Done-For-You. If you'd rather we pull five competitor libraries for you, score them, and hand you the vault — that upgrade is waiting in M-07.

DO THIS NOW • M • 00

Pick three competitors. Don't think about it too long.

Open a note. Write down three DTC brands you'd steal from if stealing were free. One should be a direct competitor. One should be a brand you admire but isn't in your category. One should be a brand you think is overrated — you want to know why their ads work anyway.

- 01 Write down the three brands. Not four. Not ten. Three.
- 02 For each, note their Facebook page name as it appears on their site (footer or social links). You'll need it in M-01.
- 03 Set a timer for five minutes. If you can't name three brands in five minutes, your problem isn't visibility — it's category awareness, and we handle that in a different course.

UP NEXT • MODULE • 01

→ p. 13

The No-Tool Path — a browser, a Claude window, and twenty minutes.

MODULE • 01 • THE SCRAPE

The No-Tool Path.

You don't need Apify. You don't need a scraper. You need a browser, a Claude window, and the willingness to let the model do what it does best — turn unstructured screens into structured rows.

READ TIME

16 minutes

ACTION AT END

Scrape one brand

TIME TO EXECUTE

Twenty minutes

§ 01 · WHY THIS PATH FIRST

Start here even if you plan to scale.

If you've never sat with a competitor's ad library and actually read it, you shouldn't automate it yet. Scraping a hundred ads into a spreadsheet is only useful if you know what to look for — and the fastest way to learn that is to do it by hand once.

The no-tool path is the training wheels that also happens to be the fastest method for operators running three brands or fewer at a time. Most of our customers never graduate past it. They don't need to. Their vault grows at the pace of their category, not at the pace of a scraper.

There are three stages. Open the Library. Capture the page. Hand Claude the screenshots. That's the whole method. Everything after this section is just doing each stage cleanly.

§ 02 · STAGE ONE · OPEN THE LIBRARY

Find the page. Filter to active. Sort by newest.

Go to facebook.com/ads/library. Set the country selector to the country you care about. If you're a US brand selling to the US, leave it on United States. If the competitor sells into Europe, flip to Germany, France, or United Kingdom — you'll get the Digital Services Act disclosures on top of the standard ones.

In the category dropdown, leave it on "All ads." The political-ads filter is separate and not what you want.

Type the brand name. Pick the right page from the dropdown — verify the page by checking the follower count, verified badge, and the page creation date against the brand's actual website. A surprising number of DTC brands have a secondary "Support" page or a regional subsidiary page running ads, and they'll pollute your capture if you grab the wrong one.

WATCHOUT · THE AGENCY PROXY

Some brands run ads out of their agency's page rather than their own. If the brand's main page shows zero active ads but you know they're spending, check for a page with "[Brand] Partner," "[Brand] Shop," or "[Brand] Official" in the name. Search the agency's page next — the ads will be there, under the agency's business manager.

Once you've got the right page loaded, flip the filter to **Active**. Sort by newest. This gives you the live roster in reverse chronological order — which is exactly how you want it, because the ads at the bottom of the page are the ones that have survived the longest, and those are the winners.

§ 03 · STAGE TWO · CAPTURE THE PAGE

Scroll, screenshot, scroll again. No tools beyond your screen recorder.

The goal here is to get the ads into Claude in a form Claude can read. You have three options, ranked by how much signal they preserve.

Option A: Long-form screenshot

On macOS use **Shift + Command + 4**, press Space, then click a Library scroll area. On Windows use the Snipping Tool's fullscreen capture. Scroll the library tab, take another screenshot at each new chunk, repeat. You'll end up with five to twenty images depending on the volume of ads the brand is running.

This is the method that works with every phone, every browser, every setup. No extensions. No logins. It's slower than the alternatives but it always works.

Option B: Full-page capture extension

If you use Chrome, install GoFullPage or a similar full-page screenshot extension. Click once, it stitches the entire scrolled page into one image. Claude reads it in a single upload. Fastest method for operators who take their vault seriously enough to spend three minutes on tooling.

Option C: Print to PDF

Chrome's print dialog, destination "Save as PDF," margins set to None. You get a multi-page PDF with the entire library captured. Claude ingests PDFs directly. This is the cleanest capture of the three and the one we use internally for batch scrapes.

Whichever option you pick, the rule is the same: **capture every active ad in one pass**. Don't stop halfway, don't spread it across two sessions. Claude needs the full slice to score the duration ladder correctly.

§ 04 · STAGE THREE · HAND IT TO CLAUDE

One prompt. One upload. One structured output.

This is the prompt. Paste it into a new Claude conversation, upload your screenshots or PDF, and let it run. The prompt is calibrated for this exact workflow — you don't need to improve it, you need to use it.

PROMPT · META LIBRARY SCRAPE · V1.0

You are a DTC competitive-intelligence analyst. I'm going to attach screenshots (or a PDF) of a brand's Meta Ad Library page.

Your job is to extract every distinct active ad and return them as structured rows. Do not summarize. Do not skip ads. Do not describe the creative – transcribe it.

For each ad, capture exactly these fields:

1. ad_id (library ID if visible, else index: A-01, A-02...)
2. start_date (the "Started running on" date, exact)
3. days_active (calculate from start_date to today)
4. variant_count (if "uses X different versions", note X; else 1)
5. format (image | video | carousel | collection)
6. platforms (FB, IG, Stories, Reels, Messenger, Audience Net)
7. primary_text (the main caption above the creative, full text)
8. headline (the bold line below the creative)
9. cta_button (Shop Now, Learn More, Sign Up, etc.)
10. landing_url (the link shown, even if shortened)
11. hook_line (first 6-10 words of primary_text – the hook)
12. offer (any discount, promo code, or free gift mentioned)
13. visual_notes (one line: what the viewer sees in frame 1)
14. angle_hypothesis (your guess at the angle: problem-aware, social proof, authority, outcome, curiosity, comparison, or value-stack)

Output as a CSV I can paste directly into Google Sheets.

First row is headers. No commentary before or after the CSV.

If a field is not visible, write "unknown" – never guess.

If the brand name is not visible on an ad, inherit it from context.

Brand I'm scraping: [BRAND NAME]

Country filter used: [COUNTRY]

Date of scrape: [TODAY]

PASTE ONCE. REUSE FOREVER. VERSION 1.0 · APRIL 2026.

Claude will return a CSV. Paste that CSV into your Swipe Vault (Google Sheet or Notion DB, both taught in M-03). You just did in twenty minutes what most agencies charge three thousand dollars to do on a discovery call.

§ 05 · WHAT TO CAPTURE — WHAT TO SKIP

Signal versus noise in the first pass.

The scrape is only useful if you're capturing the fields that actually drive decisions. Here's the honest cut.

CAPTURE · HIGH-SIGNAL FIELDS

These drive every downstream decision

START DATE	Every row. No exceptions. This is your spend proxy.
VARIANT COUNT	Maturity of testing. Winners have 3–15 variants live.
HOOK LINE	First ten words. The hook is the entire ad's job in one line.
LANDING URL	Half the signal lives on the page, not the ad. Follow every one.
OFFER	Discount, promo code, gift. Tells you price sensitivity of their target.
ANGLE	Your hypothesis. You'll refine it in M-04 scoring. Capture raw.

SKIP · LOW-SIGNAL FIELDS

Noise you'll be tempted to capture but shouldn't

EXACT RUNTIME	Start date is enough. Runtime-to-the-hour is theater.
DESCRIBE CREATIVE	Transcribe the copy, note the visual. Don't write essays.
AESTHETIC OPINIONS	Don't note that something is "clean" or "cluttered." Subjective, useless.
EVERY PLATFORM	"FB + IG" is enough. Don't note Messenger if the ad also runs on Messenger.
AGENCY PARTNER TAG	Occasionally useful in M-05 funnel audit. Skip on first pass.
FOLLOWERS, VERIFICATION	Page context is a one-time capture. Don't repeat it per ad.

§ 06 · WHERE THE NO-TOOL PATH BREAKS

Know your limits. Graduate deliberately.

This method works up to about thirty ads per brand. Past that, Claude's output quality starts degrading on single-prompt runs — you'll see it compress rows, skip the middle of the library, or start hallucinating dates. It's rare, but it happens.

If a brand has more than thirty active ads at once, split the capture: top half of the library in one Claude conversation, bottom half in another, then merge the CSVs. That buys you up to seventy-ish ads per brand without tool help.

Past seventy ads per brand — or if you're scraping more than five brands a week — you've outgrown this path. Move to M·02 and the Apify scrape. It's \$5 per run and takes the scrape time from twenty minutes to two. We'll cover setup end-to-end.

Most operators never hit this wall. The no-tool path is enough for a \$15M brand watching four direct competitors. Don't add tooling for tooling's sake.

DO THIS NOW • M • 01

Scrape competitor one. Full capture to CSV.

Pick the first brand from the list you wrote down in M·00. Run the three stages end-to-end. Do not stop to organize, clean, or score — that's M·03 and M·04. Just scrape.

- 01 Open the brand's Meta Ad Library. Filter to active. Sort by newest.
- 02 Capture the full page (long screenshot, GoFullPage, or Print-to-PDF).
- 03 Paste the v1.0 prompt into Claude. Upload your capture. Run.
- 04 Open a blank Google Sheet. Paste the CSV. Save as "[Brand] — Library Scrape — [Date]".
- 05 Do not score yet. Do not analyze yet. You're just proving the method works.

UP NEXT • MODULE • 02

→ p. 21

The Apify Path — when five brands isn't enough.

MODULE • 02 • SCALE THE SCRAPE

The Apify Path.

When five brands isn't enough — and you don't want to spend your Saturday screenshotting — a single Apify actor runs the whole capture for you in two minutes at a cost that would embarrass your agency.

READ TIME

14 minutes

ACTION AT END

Run one actor

COST PER BRAND

\$0.40–\$2.00

§ 01 · WHEN TO GRADUATE

Three signals you've outgrown the browser method.

The no-tool path is elegant until you're running it ten times a week. At that point you're doing data entry, not competitive intelligence — and the price of an afternoon of your time is more than the price of a year of Apify.

Here are the three signals. Any one of them — switch.

- **Volume.** You're tracking more than five brands and each has more than thirty active ads. That's 150+ ads per pass. The browser path breaks here.
- **Cadence.** You're scraping weekly or faster. Manual entry at that frequency is where operators quietly stop maintaining their vault.
- **Delegation.** You want to hand the scrape to an ops coordinator or EA. The Apify path fits a checklist. The browser path requires judgment.

If none of those three apply, stay on the no-tool path. Saving four dollars a month is not worth adding a SaaS dependency.

§ 02 · WHAT APIFY IS, IN ONE PARAGRAPH

A marketplace of scrapers, each one a one-click workflow.

Apify is a platform that hosts pre-built scrapers — they call them "actors" — for common websites. You don't write the scraper. You configure it with inputs, press Run, and download the output as JSON or CSV. For Meta Ad Library, two actors do everything we need. You pay by compute and dataset volume, which for our use case lands between forty cents and two dollars per brand per run.

Sign up at apify.com. The free tier gives you \$5 of monthly compute credits — enough to scrape roughly ten brands. For a weekly cadence on five brands, you'll be on the \$49/month plan. If that number horrifies you, stay on the no-tool path.

§ 03 · THE TWO ACTORS THAT MATTER

One for precision. One for volume. Pick the right one per job.

ACTOR ONE · PRECISION

curious_coder / facebook-ads-library-scraper

BEST FOR	Per-brand deep scrapes. Captures 40+ fields per ad including variant data.
INPUTS	Library URL or search query. Country. Active-status filter. Max results.
OUTPUT	JSON, detailed. Media URLs, start dates, variants, page context.
TYPICAL COST	\$0.40 per brand at ~100 ads. Pennies per ad.
SPEED	30–90 seconds per brand.
USE WHEN	You want the full field set. This is the one we use internally.

ACTOR TWO · VOLUME

apify / facebook-ads-scraper (official)

BEST FOR	Category sweeps. Scrape "skincare" or "sleep" across every brand running that term.
INPUTS	Keyword search. Country. Max ads per run.
OUTPUT	JSON, lighter field set. Good for discovery, not depth.
TYPICAL COST	\$1-\$2 per 1,000 ads.
SPEED	2-5 minutes per keyword sweep.
USE WHEN	You don't know which brands to track yet — discovery phase.

Ninety percent of operators only need the first actor. Start there.

§ 04 · CONFIGURING THE SCRAPER

Four inputs. None of them optional. All of them affect cost.

Step one: get the right Library URL

On facebook.com/ads/library, configure the filters the way you would for a browser scrape — pick country, select the brand page, filter to Active, sort by Newest. Copy the URL from your browser bar. It will look long and ugly, with dozens of query parameters. That's fine. The actor reads those parameters.

Step two: set max results

Default to 100 ads per brand. If the brand runs more than that, increase in increments of 50. Don't set it to 500 on a brand that only runs 20 active ads — you'll still pay for the compute to check.

Step three: set country

Match the country in the Library URL. If you want EU disclosures, set country to Germany or France explicitly. US is default.

Step four: output format

Select JSON output. CSV loses the variant-count array. You want JSON, and you'll convert to CSV in the next step with Claude.

COST DISCIPLINE · THE RULE OF FIFTY

Set max results to 50 on your first run of any brand. Check the output. If the actor returned fewer than 50 ads, the brand runs fewer than 50 — don't rerun at a higher cap. If it returned exactly 50, the brand runs more — rerun at 150. You'll save half your credits by not overshooting.

§ 05 · THE JSON YOU WANT — AND THE JSON TO IGNORE

Forty fields come back. Eight of them matter.

The actor returns a JSON file per brand. Each ad is an object with forty-plus fields. Most of them are metadata Meta stores for its own analytics and you don't care about. Here's the minimum viable field set you actually want.

KEEP · THESE GO IN THE SWIPE VAULT

The eight fields that drive every downstream decision

AD_ARCHIVE_ID Unique Meta ID. Use as primary key.

START_DATE Epoch or ISO string. Convert to days_active downstream.

SNAPSHOT.CARDS Array. Length = variant count.

SNAPSHOT.BODY.TEXT Primary text. Your hook source.

SNAPSHOT.TITLE Headline under the creative.

SNAPSHOT.CTA_TEXT Stop Now, Learn More, etc.
CTA_TYPE

SNAPSHOT.LINK_URL Landing page destination.

PUBLISHER_PLATFORM Array of platforms. FB, IG, Messenger, Reels.

IGNORE · FIELDS THAT WILL TEMPT YOU

Meta metadata you don't need on first pass

<code>COLLATION_ID</code>	Internal grouping. Variant count already covers this.
<code>ENTITY_TYPE</code>	Always the same for commercial ads. Noise.
<code>GATED_TYPE</code>	Political gate flag. Irrelevant for DTC.
<code>HIDDEN_SAFETY_DATA</code>	Literally the field name says it. Empty for commercial ads.
<code>PAGE_CATEGORIES[]</code>	Noisy auto-classification. Useless.
<code>IS_ACTIVE</code>	You filtered for this upstream. Always true in your output.

§ 06 · FROM JSON TO SWIPE VAULT

Let Claude flatten the structure. Don't do it by hand.

Once you have the JSON file from Apify, hand it to Claude with this prompt. It returns the same 14-column CSV you get from the browser path, formatted identically, so your vault stays consistent whether you scraped manually or with a tool.

PROMPT · APIFY JSON → VAULT CSV · V1.0

Attached is an Apify JSON export of Meta Ad Library data for one brand. Flatten it into a CSV with these 14 columns, one row per ad:

```
ad_id, start_date, days_active, variant_count, format,
platforms, primary_text, headline, cta_button, landing_url,
hook_line, offer, visual_notes, angle_hypothesis
```

Rules:

- days_active = (today - start_date) in days, integer
- variant_count = length of snapshot.cards array, min 1
- format = detect from snapshot: image | video | carousel | collection
- platforms = comma-separated from publisher_platform array
- hook_line = first 8-10 words of primary_text
- offer = extract any % off, free gift, promo code mentioned; else "none"
- visual_notes = one line describing frame 1 of the creative
- angle_hypothesis = one of: problem-aware, solution-aware, social-proof, authority, outcome, curiosity, comparison, value-stack
- if a field is null/missing, write "unknown"

Return CSV only. No preamble, no commentary, no markdown fencing.

Brand: [BRAND NAME]

Today's date: [YYYY-MM-DD]

SAME OUTPUT SHAPE AS M·01. VAULT STAYS CONSISTENT.

Paste the CSV into your vault. Done. Total time from Apify click to vault row: under five minutes per brand. Run five brands in a parallel batch — Apify will run them concurrently — and you've scraped your entire competitive set in one coffee.

§ 07 · WHAT THIS UNLOCKS

Ten brands a week isn't intimidating when each one takes a minute.

Once the Apify pipeline is running, the unlock is **cadence**. You can run weekly without it feeling like a project. You can add brands without anxiety. You can run a keyword sweep the week before a launch to see what your category is doing.

More importantly, the scrape becomes so cheap that you can afford to track brands you're curious about rather than just brands you're competing with. That's where the lateral pattern recognition lives — a sleep brand learning something from a pet food brand. You can't afford that research manually. You can easily afford it at fifty cents a run.

DO THIS NOW · M · 02

Run the Apify actor against the brand you scraped in M·01.

Yes — the same brand. You're comparing methods. It's a diagnostic, not a duplicate.

- 01 Sign up for Apify. Free tier is fine for this test.
- 02 Find the `curious_coder` facebook-ads-library-scraper actor. Clone it.
- 03 Paste in the Library URL you used in M·01. Set max to 50. Run.
- 04 Download the JSON. Hand it to Claude with the v1.0 flatten prompt.
- 05 Diff the output against your M·01 CSV. The Apify version should catch ads you missed by hand, and flag variants the browser view collapsed. That delta is why you'll graduate.

UP NEXT · MODULE · 03

→ p. 29

The Swipe Vault — where the scrape becomes a system.

MODULE • 03 • THE SYSTEM

The Swipe Vault.

A scrape without a schema is a screenshot graveyard. The vault is what turns your scrape into an instrument — something you can query, rank, and return to. This module hands you the schema.

READ TIME

12 minutes

ACTION AT END

Build the vault

WHERE IT LIVES

Notion or Sheets

§ 01 · WHY THE SCHEMA MATTERS

A vault that can't be filtered is a mood board.

The difference between a scored swipe vault and a Pinterest board is whether you can answer this question in ten seconds: "Show me every ad over thirty days old with a social-proof hook and a BOGO offer." If your system can't do that, it's decoration.

Every field in the schema below exists because it answers a question you'll actually ask when you're building a brief. The fields you won't use — runtime hours, aesthetic adjectives, engagement-counter guesses — aren't in the schema, because data you don't query is data you don't trust.

Pick one home for your vault and stay there. A vault split across three tools isn't a vault, it's a search. The two homes that work: Google Sheets (fastest to start, easiest to share) or Notion (best for annotation, slower to query). Both are covered below.

§ 02 · THE SCHEMA

Seventeen columns. Every one of them earns its place.

FIELD	TYPE	WHY IT EARNS ITS COLUMN
brand	text	Primary grouping. Keep the name canonical across rows.
ad_id	text (key)	Meta archive ID or your index. Prevents duplicate entry.
start_date	date	The spend proxy. Drives the ladder. Non-negotiable.
days_active	formula	TODAY() – start_date. Formula; never hand-enter.
variant_count	int	Maturity of testing. Winners sit at 3–15.
format	select	image, video, carousel, collection. One of four.
platforms	multi-select	FB, IG, Stories, Reels, Messenger, Audience Net.
hook_line	text	First 8–10 words of primary text. The job the ad does in one line.
hook_type	select	One of the eight taxonomy slots on p. 34.
primary_text	longtext	Full copy above the creative. Never lose this.
headline	text	Bold line under the creative. Often the offer.
cta	select	Shop Now, Learn More, Sign Up, Get Offer, Download.
offer	text	Discount, promo code, free gift, price point. "none" if absent.
landing_url	url	Click it during M-05. Half the signal lives there.
tier		Steal, Study, Skip, Ignore. Computed in M-04.

	(formula)	
score	int (formula)	0–17 from the M·04 rubric. Sort your vault on this.
notes	longtext	Your observations. One line max per note. Searchable.

If you're wondering where "spend" or "ROAS" fields went — re-read M·00. You don't have them. Every scraper that claims to hand you ROAS is guessing. Guessed metrics pollute a vault faster than missing ones.

§ 03 • THE HOOK TAXONOMY

Eight hooks. Every ad fits one. This is how you rank angles.

The hook type field is the one that takes judgment — the other sixteen fields are transcription or math. Use this taxonomy. Eight slots, mutually exclusive, chosen by the first fifteen words of the primary text.

01 • PROBLEM-AWARE

Names the pain in frame one.

Opens with the symptom. Assumes the viewer already feels it. No education required.

"Still waking up at 3am staring at the ceiling?"

02 • SOLUTION-AWARE

Names the category, not the pain.

Viewer knows solutions exist. Ad positions this one as the better option.

"The magnesium blend actually clinical sleep doctors prescribe."

03 • SOCIAL PROOF

Other people bought this.

Testimonial, review count, UGC-styled, founder story. Relies on similarity.

"47,000 women swear by this. Here's why."

04 • AUTHORITY

A credential does the heavy lifting.

Doctor, scientist, expert, clinical study, patent. Asymmetric trust.

"Stanford researchers just proved what sleep gurus have been saying."

05 • OUTCOME

Skip to the transformation.

Before-and-after. Timeline promise. Dramatic end-state without the process.

"In 14 days her skin didn't just clear — it changed."

06 • CURIOSITY

Pattern break. Unanswered question.

Opens a loop the viewer has to close. Works best in video, weakly in image.

"Why every peptide on Amazon is missing the one ingredient that matters."

07 • COMPARISON

Versus something the viewer knows.

Category leader, previous generation, DIY alternative. Frames the delta.

"Same ingredients as the \$200 brand. A third of the price."

08 • VALUE STACK

Offer-led. Price-led. Risk-reversed.

Leads with the deal. Works when the product is already familiar or commoditized.

"Three bottles. Free shipping. 90-day refund. \$39."

Every ad gets exactly one hook type. When an ad straddles two — and some will — pick the one that runs the first six words. The hook happens in the first six words or it doesn't happen.

§ 04 · BUILDING THE VAULT IN GOOGLE SHEETS

Fifteen minutes. No plugins. No templates to buy.

Open a new Google Sheet. Name it "[Your Company] — Swipe Vault — [Year]." Create the seventeen columns in row one in the order listed in the schema. That's it.

Three formulas to add, column by column.

FORMULAS · PASTE THESE

Three cells do all the math

DAYS_ACTIVE (COL D)	=IF(C2="", "", TODAY() - C2)
-------------------------------	------------------------------

TIER (COL O)	=IF(P2>=13,"Steal",IF(P2>=9,"Study",IF(P2>=5,"Skip","Ignore")))
---------------------	---

SCORE (COL P)	Hand-entered in M·04 or paste from the scoring prompt.
----------------------	--

Add conditional formatting on the tier column. Steal = green fill. Study = gold fill. Skip and Ignore = rust fill. Now at a glance you can see which rows in the vault are carrying the intelligence.

Freeze row one. Sort by score descending. Filter views save you from rebuilding the sort every time you paste new rows. That's the whole setup.

§ 05 · BUILDING THE VAULT IN NOTION

Better for annotations. Slower for sort. Worth it if you think in prose.

Create a new Database — Table view. Add the seventeen fields with the types listed in the schema (Notion's "formula," "select," and "multi-select" types map cleanly). Set the primary field to `ad_id`.

Create three linked views on the same database:

- **By Tier** — Board view, grouped by tier column. Visual triage of the full vault.
- **By Hook Type** — Gallery view, grouped by `hook_type`. For pattern analysis.
- **Steal Only** — Filtered table, tier = Steal, sorted by `days_active` descending. This is the view you open when building a brief.

One more thing Notion gives you that Sheets can't: every row is a page. Click into an ad's row and you have space to paste the creative, the landing page screenshot, your annotation prose, and links to the briefs that got built off that insight. For operators who work in writing, this is worth the slower filter experience.

DON'T SPLIT · THE VAULT RULE

Pick one. Don't run both Sheets and Notion "just to compare." The vault's value is in its queryability — and that decays the moment you have two versions of it. The best operator using Google Sheets will beat the operator running a split system every time.

§ 06 · LOADING YOUR FIRST FIFTY ROWS

The first load is the hardest. Do it tonight.

You've already scraped one brand in M-01 and M-02. Paste that CSV into your vault now. Then scrape the other two brands on your M-00 list, flatten their outputs with the same Claude prompt, and paste them in. You should land at 80–150 rows.

At 150 rows the vault stops feeling like an exercise and starts feeling like an instrument. You'll open it to answer questions. You'll sort it before a strategy meeting. You'll show it to your founder and watch them get quiet.

DO THIS NOW • M • 03

Build the vault. Load all three brands. One sitting.

If you split this across two sittings, the odds that it never gets finished go up by half. One afternoon. Start to end.

- 01 Create the Sheet or Notion DB. Add all seventeen columns.
- 02 Add the three formulas. Test them on row two.
- 03 Scrape brands two and three using the method from M·01 or M·02.
- 04 Flatten each scrape to CSV. Paste all three into your vault.
- 05 Sort by days_active descending. Read the top ten rows. Those are your controls — that's what the rest of the course teaches you to use.

UP NEXT • MODULE • 04

→ p. 37

The Scoring Layer — the rubric that tells you which rows to bet on.

MODULE • 04 • THE RANKING

The Scoring Layer.

A vault is a library. Scoring is the librarian who tells you which book to pick up. Three lenses, seventeen points, four tiers — and a prompt that does it automatically.

READ TIME

14 minutes

ACTION AT END

Score your vault

FRAMEWORK BASIS

LOOPER, manualized

§ 01 · WHAT THIS REPLACES

Every agency pitch scores ads on vibes. You're about to stop.

The version of this method we run inside our own pipeline is called LOOPER. It reads live Meta performance data and scores every active creative on a weighted rubric. You don't have the live data on a competitor — so we pulled LOOPER apart and rebuilt it around the fields you can see.

That's this module. The rubric below is LOOPER adapted for the Library. Same rigor. Different inputs. Same output: a scored, ranked, tiered list of the ads actually earning their slot.

Three lenses. Each scored on its own scale. Sum to a 0–17 composite, assign tier, sort the vault. Monday morning clarity in one afternoon.

§ 02 · LENS ONE · THE DURATION LADDER

How long has it lived? Score it 0 to 5.

Duration is the primary lens. It's the field Meta gave you that most directly proxies spend. An ad that's been live ninety days in a brand's account has cleared every kill threshold the media team uses. An ad that lived six days didn't clear one.

90+ days	Control winner · flagship	5 POINTS
60–89 days	Scaled winner	4 POINTS
30–59 days	Validated control	3 POINTS
14–29 days	In testing · surviving	2 POINTS
7–13 days	Fresh test	1 POINT
< 7 days	Too early to read	0 POINTS

One number. From the start_date column. Nothing subjective.

§ 03 · LENS TWO · THE VARIANT COUNT

How mature is their testing? Score it 0 to 4.

Variant count is a second-order signal. A brand with a ninety-day ad running six variants is doing something very different than a brand with a ninety-day ad and one. The first is actively milking a winner — iterating the creative to delay fatigue. The second is lucky, or lazy, or both.

High variant counts don't always mean high quality — but high variant counts on long-lived ads are always a signal. Score accordingly.

10+ variants	Mature iteration · mined hard	4 POINTS
5-9 variants	Actively iterating	3 POINTS
3-4 variants	Standard testing rhythm	2 POINTS
2 variants	Minimal testing	1 POINT
1 variant	Single creative · no testing	0 POINTS

THE THRASHER PATTERN

A brand running 50+ ads, all under 14 days old, none with variants — is thrashing. Their scoring will come out low across the board and that's correct. Don't study thrashers. Their account is telling you what not to do, and the not-doing is the value.

§ 04 • LENS THREE • THE CLARITY RUBRIC

Is the ad doing its job? Score it 0 to 8.

The first two lenses are about the brand's behavior. This one is about the ad itself. Four sub-criteria, two points each, totaling eight. Apply them one at a time.

CRITERION A • HOOK SHARPNESS (0-2)

2 points: First six words commit to one hook type cleanly. You can't confuse it with two other hooks. **1 point:** Hook is recognizable but shared with filler. **0 points:** First six words are brand-voice throat-clearing. "At [Brand], we believe..." = zero.

CRITERION B • OFFER CLARITY (0-2)

2 points: Offer is in the headline or first line. Price, promo, free gift, risk reversal. **1 point:** Offer exists but is buried below the fold or only in the landing page. **0 points:** No offer. "Learn more" doesn't count.

CRITERION C • VISUAL CLARITY (0-2)

2 points: Frame one reads without scrolling — you know the product and the claim in under a second. **1 point:** Frame one requires a second to parse. **0 points:** Lifestyle beauty shot with no product, no claim, no context. Common in brand-led accounts. Rarely survives ninety days.

CRITERION D • CTA MATCH (0-2)

2 points: CTA verb matches the ad's job. "Shop Now" on a sale ad, "Learn More" on an education ad, "Get Quote" on a long-sales-cycle ad. **1 point:** Generic CTA on a specific ad. **0 points:** CTA contradicts the ad's tone. "Sign Up" on a promo ad is a miss.

§ 05 • THE COMPOSITE

Sum the lenses. Four tiers. Sort the vault.

Duration 0-5 + Variants 0-4 + Clarity 0-8 = 0-17 total. Every ad in your vault ends up with a number. The number maps to a tier.

13-17	Steal • extract the DNA	TOP TIER
9-12	Study • angle, not creative	MID TIER
5-8	Skip • losing or unclear	LOW TIER
0-4	Ignore • thrasher noise	DEAD TIER

Steal is the tier you take to M·06 and convert to a brief. **Study** is the tier you note the angle of and keep in rotation — some of these move up when the brand iterates. **Skip** and **Ignore** you leave in the vault for pattern analysis but don't mine for direction.

A healthy vault has 15–25 percent of rows in the Steal tier, 30–40 percent in Study, and the balance in Skip or Ignore. If your Steal tier is over 40 percent, your rubric is loose — re-score with the 2-point criteria tightened.

§ 06 · SCORE IT WITH CLAUDE

You won't do this by hand after the first twenty. Automate it.

Once the rubric is in your head, hand it to Claude with your vault CSV and let it score the rest. Here's the prompt. It knows the rubric, knows the tier bands, and returns your vault with three new columns populated.

PROMPT • AUTO-SCORE THE VAULT • V1.0

You are scoring competitor ads using the Steal the Signal rubric.

Attached is my swipe vault CSV. Return it unchanged except for three columns you will fill in for every row:

```
duration_points    (0-5, based on days_active)
variant_points     (0-4, based on variant_count)
clarity_points     (0-8, based on the clarity rubric below)
```

Then add two derived columns:

```
score = duration_points + variant_points + clarity_points (0-17)
tier  = "Steal" if 13+, "Study" if 9-12, "Skip" if 5-8, "Ignore"
        if 0-4
```

Duration points:

```
90+ days = 5 · 60-89 = 4 · 30-59 = 3 · 14-29 = 2 · 7-13 = 1
< 7 = 0
```

Variant points:

```
10+ = 4 · 5-9 = 3 · 3-4 = 2 · 2 = 1 · 1 = 0
```

Clarity rubric (score 0-2 each, sum for clarity_points):

- A. Hook sharpness – first 6 words commit to one clean hook type
- B. Offer clarity – offer visible in headline or first line
- C. Visual clarity – frame 1 reads in under one second
- D. CTA match – CTA verb matches the ad's job

Be strict. Do not give 2/2 unless the criterion is obviously met. Thrasher ads (many short-lived, no variants, muddy hooks) should score 0-4 and tier as "Ignore."

Return the full CSV with the five new columns appended. No preamble. No commentary. Just CSV.

PASTE ONCE PER VAULT REFRESH. TAKES CLAUDE 30 SECONDS TO SCORE 150 ROWS.

§ 07 · READING THE SCORED VAULT

Three questions the scored vault answers in under a minute.

Question one: what's the highest-scoring ad across all your competitors?

Sort by score descending. Top row. Read the hook line. Open the landing page. Twenty minutes of focused study on that one row teaches you more about your category than three weeks of Slack chatter.

Question two: which hook type dominates the Steal tier?

Filter tier = Steal. Group by hook_type. Count rows per group. Whichever hook type has the most rows is the angle your category is currently rewarding. That's the angle your next brief should explore — not because it's trendy, because the Library just told you it pays.

Question three: where's the gap?

Which hook type shows up zero or one times in the Steal tier? That's either an angle that doesn't work for your category (fine, move on) or an angle nobody has tried yet (rare, valuable). The second case is what operators mean when they say "whitespace." Now you can see it.

DO THIS NOW • M • 04

Score the vault. Apply Claude's rubric. Read the Steal tier.

You're about to find out which three ads in your vault are the ones worth studying — and which seventy are noise.

- 01 Export your vault to CSV.
- 02 Paste it into Claude with the v1.0 scoring prompt.
- 03 Paste the returned CSV back into your vault, replacing the old one or appending the new columns.
- 04 Sort descending by score. Filter tier = Steal.
- 05 Read every row in the Steal tier end to end. Open every landing page. Take notes. That's where M·05 lives.

UP NEXT • MODULE • 05

→ p. 45

Steal the Funnel, Not the Ad — why the landing page is half the signal.

MODULE • 05 • THE FUNNEL AUDIT

Steal the Funnel, Not the Ad.

A winning ad is half a winning funnel. Most mid-market brands copy the creative, skip the page, and wonder why their ROAS doesn't move. The ad is not the product. The page is the product — and that's where most competitors bleed.

READ TIME

14 minutes

ACTION AT END

Audit three pages

FRAMEWORK

Seven-point match

§ 01 · THE HALF NOBODY AUDITS

The landing page carries the ad, not the other way around.

An ad's job is to earn a click. A page's job is to earn a purchase. A long-running ad on top of a mismatched page is a brand that's burning CAC without knowing why — and the Library will show you the ad all day long, but it will never show you the page unless you click.

The operators who actually beat their competitors in DTC aren't the ones with the best creative. They're the ones running a B+ ad on an A+ page while their competitor runs an A+ ad on a C page. The page is the compounding moat. The ad is the bait.

This module is the audit method. Seven checkpoints between the ad and the page. A winning ad passes all seven. A losing ad misses at least two — and if you're about to copy a competitor's creative and ship it against your own page, the audit will tell you before you waste the budget.

§ 02 · HOW TO AUDIT A PAGE

Click the link. Take two screenshots. Answer seven questions.

The workflow is mechanical. For every ad in the Steal tier of your vault, click the landing_url column. Screenshot the first viewport (hero) and the first scroll (second section). That's all you need for the audit.

Open your notes. Run the seven-point check against the ad and the page side by side. Every point is a yes or no. Tally the yes count at the end.

OPERATOR DISCIPLINE · DO THIS LIVE

Don't save the page audits for later. The moment you finish scoring the vault in M-04, open the Steal-tier rows one by one and audit the page right then. It takes three minutes per ad. An audited row is an actionable row. An unaudited row is a dead row.

§ 03 · THE SEVEN-POINT MATCH

One checklist. Seven passes. Apply it to every Steal-tier ad.

THE AD-TO-PAGE AUDIT · V1.0

- ❑ 1. Hook continuity. The first six words of the ad reappear — nearly verbatim — somewhere in the page's hero section. The viewer's working memory is a five-second window. Break it and they bounce.

- ❑ 2. Offer match. The offer promised in the ad (discount, free gift, promo code, price anchor) is visible above the fold on the page, with identical language. A 20% off ad on a page that only shows the MSRP is a broken funnel.

- ❑ 3. Visual continuity. The product or scene from the ad's hero frame appears in the page's hero image. Same lighting, same styling, ideally the same shot. Different photography tells the viewer they clicked a trick.

- ❑ 4. CTA congruence. "Shop Now" on the ad, "Shop Now" or "Buy Now" on the page. "Learn More" on the ad leads to a content page, not a PDP. A content-ad-to-PDP is one of the most common funnel errors in DTC.

- ❑ 5. Proof parity. The social proof signal used in the ad is re-used on the page. If the ad says "47,000 customers," the page shows review counts or testimonials at that order of magnitude. Number mismatch kills trust.

- ❑ 6. Speed to commit. Add-to-cart or checkout is visible within one viewport on mobile. No pop-up between ad click and cart button. No email capture before purchase is possible. Every gate costs conversions.

- ❑ 7. Price anchoring. The price is visible above the fold. The comparison price (if any) is formatted to read as a saving, not a sticker shock. If the ad sells a \$39 product and the page shows \$149 with a discount code not yet applied, the math is broken.

Seven yeses is rare — we see it maybe once in every twenty-five Steal-tier ads. Five or six yeses is common for brands that have worked their funnel. Three or fewer yeses is a brand running an

ad on top of a page that doesn't match — they're getting lucky on volume, not on funnel construction.

§ 04 · WHAT TO DO WITH THE AUDIT

Three different scenarios. Three different plays.

Scenario A: 6–7 yeses

You're looking at a complete funnel. This is the one to study hardest. The competitor isn't just running a creative that works — they've built the whole machine behind it. Extract both the ad DNA and the page DNA in your M-06 brief. Ignore one and you're back to shipping a B+ ad onto your own B- page.

Scenario B: 3–5 yeses

Common. The competitor has a working ad on a partially broken funnel. The opportunity is to copy their angle and out-build them on the page. If their ad scored 16/17 in M-04 but their page scores 3/7 on the match audit — that gap is your edge. Same hook, better page, better conversion. This is the most common winnable fight in the DTC mid-market.

Scenario C: 0–2 yeses

The ad works despite the page, not because of it. Either the brand has a cult audience that converts on intent alone, or the ad is subsidizing a broken machine with heroic creative. Rare. Usually found in brands with strong organic presence or a founder with an audience. Steal the ad but don't study the funnel — there isn't one.

§ 05 · THE COMMON BLEEDS

Six patterns you'll see repeatedly. Name them fast.

THE SIX FUNNEL BLEEDS

Once you can name them, you stop making them

CONTENT-TO-PDP	"Learn more" style educational ad that drops directly onto a product page. The viewer came for a story and got a shelf.
OFFER DRIFT	Ad promises 25% off, page shows 20%. Happens when promos get updated in the PIM but not in the ad account.
HERO SWAP	Ad hero is one product, page hero is a lifestyle shot without that product. Visual whiplash. Kills trust.
GATE BEFORE CART	Email capture pop-up or age gate fires before the viewer has seen the price. Friction before desire.
REVIEW COUNT MISMATCH	Ad says "thousands of reviews," page shows 127. Either the ad's lying or the reviews are hidden below the fold.
MOBILE-DESKTOP DESYNC	Ad was designed for mobile but page tested on desktop. Hero image crops the product off-screen on a 390px viewport.

Every one of these is fixable in a sprint. Half of them are fixable in an hour. The fact that competitors have them live is not an accident — it's the compounding ground truth that most DTC brands never audit the seam between their ad account and their Shopify stack.

§ 06 · SCORING THE PAGE INTO THE VAULT

Add two columns. Keep the audit live.

Add two columns to your vault schema: **funnel_yes** (count of yeses from the 7-point audit) and **funnel_notes** (one line per audited ad, your observations). Now every Steal-tier row has a paired funnel score, and you can sort the vault by the combined (score + funnel_yes) to find the complete winners — the ads that have both a great hook and a great page.

If you want Claude to handle it, paste the ad-row plus the two page screenshots into a chat with this prompt:

PROMPT · FUNNEL AUDIT · V1.0

Attached: one ad (copy + headline + offer) and two landing-page screenshots (hero + second scroll). Apply the seven-point ad-to-page audit from Steal the Signal M05. For each of the seven criteria, return YES or NO, plus a one-line justification (< 12 words). Then return the total yes count.

The seven criteria:

1. Hook continuity — first 6 words of ad appear in hero
2. Offer match — ad offer visible above the page fold
3. Visual continuity — ad hero image continues on the page
4. CTA congruence — ad CTA verb matches the page CTA
5. Proof parity — ad social-proof signal present on page
6. Speed to commit — cart/checkout within one mobile viewport
7. Price anchoring — price visible, comparison reads as savings

Return as a 7-row markdown table. No preamble.

THREE MINUTES PER AD. SCALES TO THE WHOLE STEAL TIER.

DO THIS NOW • M • 05

Audit the top three ads in your Steal tier. Pages and all.

This is the work that changes how you see your own funnel. You'll catch your own bleeds inside of a week.

- 01 Open the top three rows of your Steal tier.
- 02 Click every landing_url. Screenshot hero and second scroll.
- 03 Run the seven-point audit by hand or with the v1.0 prompt.
- 04 Log funnel_yes and funnel_notes into your vault.
- 05 After the third audit, flip the browser to your own landing pages and run the same audit. What you learn about yourself in that ten minutes will pay for this field guide by itself.

UP NEXT • MODULE • 06

→ p. 53

Signal to Brief — converting the vault into creative direction.

MODULE • 06 • CONVERT TO DIRECTION

Signal to Brief.

A scored vault is a report. A brief is a decision. This module is the translation layer — how to take one row of stolen signal and turn it into creative direction your team, studio, or AI pipeline can execute without asking a follow-up question.

READ TIME

13 minutes

ACTION AT END

Write one brief

TEMPLATE

12 fields, fill-in

§ 01 • THE TRAP YOU'RE ABOUT TO WALK AROUND

Stop copying the creative. Steal the structure.

When operators first get a scored vault, the instinct is to save the JPGs and hand them to a creative team with a note that says "make this." That's the trap. What you want to steal is not the aesthetic — it's the structural DNA of why the ad works.

Structural DNA is the five-part spine underneath every winning ad: the hook type, the angle, the proof mechanism, the offer construction, and the format. Replicate those five and you'll produce creative that works in your brand's voice. Copy the art and you'll produce a flat imitation that reads as imitation to the algorithm and to the customer.

The difference sounds academic until you see it work. A peptide brand we worked with studied a longevity competitor's winning ad — authority hook, white-coat format, outcome-driven proof stack, money-back offer. They didn't copy the creative. They took the five-part DNA and rebuilt it in their own voice, around their own founder, with their own ingredient story. The result went live for sixty days. Theirs did too.

§ 02 · THE FIVE-PART DNA EXTRACTION

Every winning ad breaks into these five questions.

STRUCTURAL DNA · FIVE-PART EXTRACTION

Answer each question about the source ad. Don't describe the creative.

HOOK TYPE	Which of the 8 from M-03? Be specific — if it straddles two, pick the one in the first six words.
ANGLE	The one-sentence argument the ad is making. "Taking this fixes that." Not the product claim — the underlying logic.
PROOF MECHANISM	How does the ad prove the claim? Testimonial, clinical, demo, transformation, authority, UGC, data viz.
OFFER CONSTRUCTION	Price + promo + risk-reversal + urgency. How is the money risk handled?
FORMAT	Video length, image composition, talking head vs. b-roll, UGC vs. studio. The container the DNA sits in.

Those five, written out in prose for one competitor ad, take you eight minutes. That is the extract. That is what goes into the brief. The JPG doesn't.

§ 03 • THE BRIEF TEMPLATE

Twelve fields. Fill every one. Never skip to execution.

Below is the template. It's calibrated to feed either a human creative team or an AI generation pipeline — we run it through our own FORGE pipeline inside the BIS Dashboard, and we've tested it against three different creative studios. Twelve fields. Every one earns its slot.

TEMPLATE • STRUCTURAL DNA BRIEF • V1.0

Steal the Signal — Ad Brief

01 • SOURCE

Which competitor ad did this DNA come from? Brand + ad_id + days_active.

02 • HOOK TYPE

One of the eight from M-03.

03 • ANGLE (ONE SENTENCE)

The argument the ad is making. Translate into your brand's voice before writing.

04 • PROOF MECHANISM

How you'll prove the angle — mechanism must exist in your product or founder story.

05 • OFFER CONSTRUCTION

Price, promo, risk reversal, urgency. What's your version of their deal?

06 • FORMAT

Video length, aspect, talking head or b-roll, studio or UGC.

07 • FIRST SIX WORDS

The literal hook line. Write three versions. Pick the sharpest.

08 • VISUAL FRAME ONE

Single sentence describing the first frame the viewer sees.

09 • CTA VERB

Shop Now, Learn More, Sign Up, Get Offer. Must match the landing page's CTA.

10 • LANDING PAGE REQUIREMENT

Which of the seven funnel-match criteria must the page satisfy? (M-05)

11 • WHAT TO AVOID

What would make this a copy instead of a remix? Name two to three specifics.

12 • SUCCESS MARKER

How you'll know in 14 days whether the ad is a candidate to scale.

§ 04 • AUTO-FILL WITH CLAUDE

Once you've scored a vault, Claude can draft 80 percent of the brief.

The prompt below takes one row of your scored vault plus the Funnel Audit output from M-05 and returns the brief template populated. You review and sharpen — don't skip this step, the brief is the single highest-leverage deliverable in the course — but Claude does the typing.

PROMPT • VAULT ROW → BRIEF • V1.0

Attached is one row from my Steal the Signal swipe vault (a scored competitor ad, tier "Steal") plus my notes from the M-05 funnel audit for that row.

My brand: [YOUR BRAND, one sentence of what it sells]

My founder story / core proof asset: [ONE SENTENCE]

My landing page URL: [URL]

Draft a filled-in 12-field brief using the Structural DNA Brief template from Steal the Signal M-06. Your output should:

- Extract the 5-part DNA from the source ad honestly
- Translate the angle into my brand's voice and proof asset
- Write three versions of field 7 (first-six-words hook)
- Flag any field where the competitor's funnel audit scored NO and instruct how to fix it on my side
- Not copy any literal phrase from the source ad — remix the structure, rewrite the language

Output as a numbered list matching the 12 fields. Where a field has multiple options (e.g. field 7), label them A/B/C. End with a one-line "Success marker at day 14."

REVIEW, SHARPEN, SHIP. CLAUDE DOES THE TYPING; YOU DO THE THINKING.

§ 05 · SHIPPING THE BRIEF

Who gets it, and how fast they should move.

If you have an in-house creative team, the brief goes to them with the scored vault row and the funnel-audit screenshots attached. Their job is to produce three creative concepts within a week, one per version of field 7. No "we need more direction" Slacks — the brief is the direction.

If you have an external studio, the brief goes to them with a one-week deadline. Studios that say they need three weeks to produce three concepts off a twelve-field brief are overpriced. This isn't a Super Bowl spot.

If you run an AI pipeline (our own FORGE pipeline, or a comparable Midjourney + ElevenLabs + Runway stack), the brief fields map almost one-to-one into the prompt chain. Field 8 is the Midjourney frame-one prompt. Field 7 is the VO script opener. Field 5 becomes the CTA overlay. Fields 10 and 11 become the negative prompts. Three creative concepts in an afternoon, not a sprint.

WEEKLY RHYTHM · THE COMPOUNDING BIT

One brief a week, derived from one fresh vault scrape a week, produced in one executorial cycle per week. Four weeks = four briefs = twelve creative concepts tested. This is the rhythm that breaks the plateau. Most operators do this exactly once. The compounding belongs to the ones who do it every Monday.

DO THIS NOW • M • 06

Write one brief. The top row of your Steal tier.

Not three briefs. Not ten. One. Do it all the way through the twelve fields — no blanks, no "to be determined," no skipping field 11.

- 01 Open the top row of your Steal tier and the funnel audit from M-05.
- 02 Paste both into Claude with the v1.0 brief prompt.
- 03 Read Claude's draft end to end. Sharpen every field in your own words.
- 04 Send it to whoever ships your creative — in-house, studio, or pipeline.
- 05 Set a 14-day calendar reminder titled "Check M-06 success marker." That's your decision gate.

UP NEXT • MODULE • 07

→ p. 61

What's Next — the rhythm, the DFY, and the upgrade paths.

MODULE · 07 · THE COMPOUNDING

What's Next.

Every operator who finishes the course gets to the same crossroads. Run the rhythm yourself, let us run the first five libraries for you, or graduate into the automated version that runs every week without you touching it.

READ TIME

10 minutes

ACTION AT END

Pick your path

THREE PATHS

DIY / DFY / Auto

§ 01 · THE WEEKLY RHYTHM

Forty-five minutes. Once a week. That's the whole game.

The operators who get compounding value out of this field guide don't do anything fancy. They put a recurring event on the calendar. Monday, 9:00am, forty-five minutes, every week. The rhythm is the product.

Minutes 0-10	Re-scrape your three brands	APIFY OR BROWSER
Minutes 10-20	Paste new rows, let Claude score	V1.0 SCORING PROMPT
Minutes 20-30	Read what changed	NEW STEALS, FALLEN STUDYS
Minutes 30-45	Write one brief — maybe	SKIP WEEKS THE VAULT IS QUIET

That's the whole thing. Four quarters. Same four every week. You don't need to write a brief every week — most weeks, your current brief is still in execution. But you need to see the vault every week. What's new, what just hit 30 days, what fell out of the Steal tier, what brand you weren't tracking just showed up in a keyword sweep.

The forty-five minutes is what separates operators who break the plateau from the ones who stay on it. Not the scrape. Not the score. The rhythm.

§ 02 · THE FIVE-LIBRARY DONE-FOR-YOU

If the rhythm is going to slip, let us run the first quarter for you.

Most operators reading this are going to go home, run M-00 and M-01, and then life will happen. Q4 launches. Investor meeting. Ops fire. The vault goes cold. The rhythm never starts.

For those operators, we run a Done-For-You option. You name five competitors. We scrape the full library on each. We score every ad. We run the funnel audit on every Steal-tier row. We hand you a populated vault, three pre-written briefs, and a thirty-minute call to walk you through the top-of-Steal plays in your category.

DONE-FOR-YOU · FIVE LIBRARIES

We pull five competitor libraries. You get a scored vault and three briefs.

Seven days from intake to delivery. Populated Swipe Vault, 150+ scored rows, three pre-written Structural DNA briefs, thirty-minute handover call. Designed for the operator who knows they won't run the rhythm themselves, but knows the vault would change how they see their account.

DELIVERY · 7 DAYS · INCLUDES 30-MIN CALL · ADD-ON
TO THIS FIELD GUIDE

§ 03 · THE UPGRADE PATHS

Three tiers, one direction. Start wherever you are.

This field guide is the \$97 entry point. It teaches the manual method end to end. Most operators run it for a quarter, get the rhythm working, and then ask the same question: "Is there a version of this that runs itself?"

There are two. Here's the full stack.

ENTRY · THE FIELD GUIDE Steal the Signal	The book you're holding. Manual method. Eight modules, the kit, the prompts, the DFY option. Self-paced.	\$97
MOST OPERATORS · THE COURSE The Brand Intelligence System	The full course. Three modules that take Steal the Signal and extend it into pricing intelligence, ICP extraction, and message architecture. Manual method across five platforms. Video + kit + community.	\$997
SCALED OPERATORS · SAAS BIS Dashboard · Monthly	The manual method automated. 120 Claude-powered prompts, API-connected to Meta, Instagram, TikTok, Reddit, and Google Trends. Weekly creative iteration output. Continuous intelligence feed.	\$249 per month
BEST VALUE · ANNUAL BIS Dashboard · Annual	Same dashboard. Pay annually. Saves \$994 vs the monthly plan. Includes quarterly brief review call.	\$1,997 per year
FULL STACK · BUNDLE Course + Dashboard · Annual	The course teaches it. The Dashboard runs it. Everything, together. Save \$497 vs buying them separately.	\$2,497

No upgrade pressure. Most operators stay on the \$97 field guide for at least a quarter. The people who move up are the ones who've already proven the rhythm — and want to give a tool back to their team so they don't have to personally run it.

§ 04 · WHEN TO MOVE UP A TIER

Three honest signals you've outgrown the field guide.

- **Signal one · You're running the rhythm.** You've run the Monday scrape four weeks in a row. You've written four briefs. You know the method and you want to stop typing. Move to the Dashboard.
- **Signal two · You're the bottleneck.** You want your team to run this but they can't execute a CSV scrape on their own. Move to the Course — the video walkthrough is the training module for your team.
- **Signal three · You want more than ads.** The Library gives you ads. But pricing, ICP extraction, message architecture, and competitor positioning live on five other platforms — Reddit, Google Trends, TikTok search, Instagram, and the brand's own site. BIS (\$997) is the method across all five. The Dashboard is the automation across all five.

§ 05 · THE CLOSER

You don't have a marketing problem. You have a visibility problem.

That was the line in M-00. It's the line the whole field guide proves. Every module after the reframe was a method for closing the visibility gap — on your competitors, on your own funnel, on the angle your category is currently rewarding.

You can see it now. You can't unsee it. You can't go back to scrolling the Ad Library for an hour and walking away with nothing. The method changed your relationship to the data, and the vault you built in M-03 is the compounding instrument.

Most DTC brands hit a wall and never figure out why. You just built the instrument that tells you why. What you do with that next — run it Monday mornings, hand it to us, or graduate into the automated version — is the only question left.

NEXT STEP · BRAND INTELLIGENCE SYSTEM

Give me the system — \$997.

The full course. Three modules that take Steal the Signal and extend it across every surface where your customers speak. Pricing intelligence. ICP extraction. Message architecture. The system that makes you the most valuable person in the room — and everyone in the room knows it.

SELF-PACED · LIFETIME ACCESS · DASHBOARD
UPGRADE AVAILABLE

UP NEXT · APPENDIX

→ p. 69

The Kit — every prompt, template, and checklist in one reference.

APPENDIX • THE KIT

Every prompt, templated.

A single-reference index of the prompts, checklists, and templates distributed across the course. Use this page to find anything fast during a live run. All assets also ship as standalone files.

PROMPTS

5 operator-ready

TEMPLATES

Vault • Brief • Audit

FORMAT

Copy, paste, ship

A · THE PROMPT LIBRARY

Five prompts. Numbered. Version-locked.

Each prompt is reproduced in full in its home module. This index points you to where each lives. If you treat the prompts as living documents — tweak, save, re-test — version them with your own suffix, not ours.

P · 01	Meta Library Scrape Screenshots or PDF of a library page → 14-column CSV. Home in M·01.	V1.0 · M·01
P · 02	Apify JSON » Vault CSV Flattens Apify raw JSON into the same 14-column schema. Home in M·02.	V1.0 · M·02
P · 03	Auto-Score the Vault Scores every row on duration + variants + clarity. Returns tier. Home in M·04.	V1.0 · M·04
P · 04	Funnel Audit Seven-point ad-to-page match. Runs against screenshots. Home in M·05.	V1.0 · M·05
P · 05	Vault Row » Brief Converts one Steal-tier row into a 12-field Structural DNA Brief. Home in M·06.	V1.0 · M·06

B · VAULT SCHEMA · ONE-PAGE REFERENCE

The columns, the formulas, the tier bands.

COLUMN	TYPE	SOURCE / FORMULA
brand	text	Hand-enter, canonical.
ad_id	text (key)	Meta archive ID.
start_date	date	From scrape.
days_active	formula	=TODAY()-start_date
variant_count	int	From scrape.
format	select	image / video / carousel / collection
platforms	multi-select	FB, IG, Stories, Reels, Messenger
hook_line	text	First 8–10 words of primary_text.
hook_type	select	One of the 8 taxonomy slots.
primary_text	longtext	Full ad copy.
headline	text	From scrape.
cta	select	Shop Now, Learn More, Sign Up, Get Offer.
offer	text	Discount, promo, gift. "none" if absent.
landing_url	url	From scrape.
score	int (0-17)	Sum of duration + variant + clarity.
tier	formula	=IF(score>=13,"Steal",IF(>=9,"Study",IF(>=5,"Skip","Igr
funnel_yes	int (0-7)	From M-05 audit.

notes

longtext

One line per note.

C · TIER BANDS · POCKET REFERENCE

13–17	Steal · extract DNA, write brief	TOP
9–12	Study · note angle, watch weekly	MID
5–8	Skip · losing, keep for context	LOW
0–4	Ignore · thrasher noise	DEAD

D · THE HOOK TAXONOMY · EIGHT SLOTS

- **01 Problem-Aware** — names the pain in the first six words.
- **02 Solution-Aware** — names the category, positions as the better option.
- **03 Social Proof** — other people bought this. Testimonial, review count, UGC.
- **04 Authority** — credential, doctor, study, patent. Asymmetric trust.
- **05 Outcome** — skip to the transformation. Before-and-after, timeline.
- **06 Curiosity** — pattern break, unanswered question, loop.
- **07 Comparison** — versus category leader, previous gen, DIY alternative.
- **08 Value Stack** — offer-led, price-led, risk-reversed.

E · THE SEVEN-POINT FUNNEL AUDIT

APPLY TO EVERY STEAL-TIER ROW

- ☐ Hook continuity — ad's first 6 words reappear in page hero.

- ☐ Offer match — promised offer visible above fold on the page.

- ☐ Visual continuity — ad hero image continues on the page.

- ☐ CTA congruence — ad CTA verb = page CTA verb.

- ☐ Proof parity — social proof signal survives the click.

- ☐ Speed to commit — cart/checkout within one mobile viewport.

- ☐ Price anchoring — price visible, comparison reads as savings.

F · THE WEEKLY RHYTHM

0-10 min	Re-scrape 3 brands	APIFY PREFERRED
10-20 min	Claude scores new rows	PROMPT P · 03
20-30 min	Read deltas — new Steals, fallen Studys	QUIET = SKIP
30-45 min	One brief — only if new Steal appeared	PROMPT P · 05

G · A FINAL NOTE

Every tool in this book is one you'll use once and replace with a better version of your own. That's fine. Version them. Ship. What matters is that you close the gap between "I can see what they're doing" and "I'm using it on Monday." That's the whole game.

END OF FIELD GUIDE · OPERATOR EDITION 01 · IRREVERENT MARKETING ·
IRREV.CO/BIS